

Machine/Statistical Learning  
Introduction to Neural Networks,  
Back propagation principle

Florent Chatelain, Olivier.JJ.Michel


Grenoble-INP, GIPSA-Lab

ENSTA, feb. 2020

## Motivation

- Build a Bio-inspired **parametric** model, with possibly high complexity.

## Rosenblatt's perceptron, 1957



An illustration of an artificial neuron. Source: Becoming Human.

- ▶ Neural like structure, with a single unit
- ▶  $n$  inputs

$w_j \stackrel{def}{=} \text{weight, or connecting weight}$

Let  $g(\cdot)$  be the **activation** function, and  $x = [x_1, \dots, x_n]$  the input :

$$y = g(a(x))$$
$$a(x) = \sum_{j=1}^n w_j x_j$$

## Remarks

- ▶ -for a (linear) classification problem,  $g(a)$  is a threshold (or sigmoidal) function  
for a (linear) regression problem,  $g(a) = a$
- ▶ for binary classification,

$$g(a) = \begin{cases} -1 & \text{if } a \leq 0 \\ 1 & \text{if } a > 0 \end{cases}$$

- ▶ It is often convenient to introduce a bias to account for possible affine separating hyperplane :  
$$\begin{aligned} [x_1, \dots, x_n] &\leftarrow [1, x_1, \dots, x_n] = x \\ [w_1, \dots, w_n] &\leftarrow [w_0, w_1, \dots, w_n] = w \end{aligned}$$
- ▶ In order to predict the probability of  $x$  to be in a given class :

$$g(a(x)) = \frac{1}{1 + e^{a(x)}}$$

See `MLPC_logical_example.ipynb`

## Training the perceptron

As for other ML approaches, minimize the empirical risk, i.e. an averaged cost function.

### Online learning


The weights (parameters)  $w_0, \dots, w_n$  are updated to minimize the risk  $L(f(x^i), y^i)$  each time that a new pair  $(x^i, y^i)$  is received, as **opposed to batch learning**. Gradient descent algorithm minimization :

$$w_j \leftarrow w_j - \nu \frac{\partial L(f(x^i), y^i)}{\partial w_j}$$

- ▶  $\nu$  is the **learning rate**. In practice,  $\nu$  is often decreased when the risk is close to the minimum.
  - ▶ if  $\nu$  is too large : possible instability
  - ▶ if  $\nu$  is too small : slow convergence
- ▶ Many **epochs** may be performed on the whole training set.

## historical example : binary classifier

- ▶  $y^i \in \{-1, +1\} \forall i \in [1, N]$
- ▶  $L(f(x^i), y^i) = \max(0, -y^i a(x^i)) = \max(0, -y^i w^T x^i)$



$$\Rightarrow w_j \leftarrow \begin{cases} w_j & \text{if } y^i w^T x^i > 0 \\ w_j - \nu y^i w^T x^i & \text{if } y^i w^T x^i \leq 0 \end{cases}$$

Albert Novikov theorem, 1962

Let  $\mathcal{T} = \{(x^i, y^i), i = 1 \dots, N\}$  be the training set. Let  $D, \gamma \in \mathbb{R}^{+*}$ , then

**IF**


- $\forall x^i \in \mathcal{T}, \|x^i\|^2 < D$  ( $\leftarrow$  bounded support)
- $\exists u \in \mathbb{R}^{n+1} / \|u\|^2 = 1$  and  $\forall (x^i, y^i) \in \mathcal{T}, y^i u^T x^i \geq \gamma$  ( $\leftarrow$  margin condition)

**THEN** the perceptron algorithm converges in less than  $\left(\frac{D}{\gamma}\right)^2$  iterations.

See [Perceptron\\_sonar\\_example.ipynb](#)

## Motivation

Allow to deal with non linear frontiers between classes by inserting **hidden** layers, within a **FEED FORWARD** network.



## Notations

$w_{ij}^k$  : weight for node  $j$  in layer  $l_k$  for incoming node  $i$

$b_i^k$  : bias for node  $i$  in layer  $l_k$

$a_i^k$  : product sum plus bias (activation) for node  $i$  in layer  $l_k$

$o_i^k$  : output for node  $i$  in layer  $l_k$


$r_k$  : number of nodes in layer  $l_k$

$g$  : activation function for the hidden layer nodes

$g_o$  : activation function for the output layer nodes

Then

$$a_j^k = b_j^k + \sum_{i=1}^{r_{k-1}} w_{ij}^k o_i^{k-1} = \sum_{i=0}^{r_{k-1}} w_{ij}^k o_i^{k-1}$$



$$o_j^k = g(a_j^k) = g(\sum_{i=0}^{r_{k-1}} w_{ij}^k o_i^{k-1})$$

## Gradient backpropagation algorithm

Backpropagation attempts to minimize the empirical risk (or loss)

$$L(X, \theta) = \frac{1}{N} \sum_i L(f(x^i), y^i)$$

with respect to the neural network's weights (gathered in  $\theta$ ) :

→ for each weight  $w_{ij}^k$ , evaluate  $\frac{\partial L}{\partial w_{ij}^k}$ . By decomposing into a sum over individual error terms for each individual input-output pair

$$\frac{\partial L(X, \theta)}{\partial w_{ij}^k} = \frac{1}{N} \sum_{d=1}^N \frac{\partial L(f(x^d), y^d)}{\partial w_{ij}^k} = \frac{1}{N} \sum_{d=1}^N \frac{\partial L_d}{\partial w_{ij}^k}.$$



## Loss function derivatives (index $d$ is omitted hereafter)

Remind the expression of output at node  $N_j$  :

$$o_j^k = g(a_j^k) = g\left(\sum_{i=0}^{r_{k-1}} w_{ij}^k o_i^{k-1}\right)$$

$L$  depends on  $w_{ij}^k$  through  $a_j^k$  : Apply the chain rule to the loss function partial derivative

$$\frac{\partial L}{\partial w_{ij}^k} = \frac{\partial L}{\partial a_j^k} \frac{\partial a_j^k}{\partial w_{ij}^k}$$

$$\left\{ \begin{array}{l} \delta_j^k \stackrel{def}{=} \frac{\partial L}{\partial a_j^k} \\ \frac{\partial a_j^k}{\partial w_{ij}^k} = \frac{\partial}{\partial w_{ij}^k} \left( \sum_{l=0}^{r_{k-1}} w_{lj}^k o_l^{k-1} \right) = o_i^{k-1} \end{array} \right. \Rightarrow \frac{\partial L}{\partial w_{ij}^k} = \delta_j^k o_i^{k-1}.$$

## Output Layer (MLP with $m + 1$ layers)

Assume a one-output neural network, so there is only one output node ( $j = 1$ ). Expressing  $L$  in terms of the value  $a_1^m$  (since  $\delta_1^m$  is a partial derivative with respect to  $a_1^m$ ) gives

$$L(f(x), y) = L(g_o(a_1^m), y)$$

where  $g_o(x)$  is the activation function for the output layer.  
thus,

$$\delta_1^m = L' (g_o(a_1^m), y) g'_o(a_1^m)$$

and finally

$$\frac{\partial L}{\partial w_{i1}^m} = \delta_1^m o_i^{m-1} = L' (g_o(a_1^m), y) g'_o(a_1^m) o_i^{m-1}.$$

## Backpropagation for hidden layers -cont'd- :

One has for the error term  $\delta_j^k$  in layer  $1 \leq k < m$  :

$$\delta_j^k = \frac{\partial L}{\partial a_j^k} = \sum_{l=1}^{r^{k+1}} \frac{\partial L}{\partial a_l^{k+1}} \frac{\partial a_l^{k+1}}{\partial a_j^k},$$

where  $l$  ranges from 1 to  $r^{k+1}$

(The bias input  $a_0^k$  corresponds to  $w_{0j}^{k+1}$  is fixed, does not depend on the outputs of previous layers, thus  $l$  does not take on the value 0.)

Plugging in the error term :

$$\delta_j^k = \sum_{l=1}^{r^{k+1}} \delta_l^{k+1} \frac{\partial a_l^{k+1}}{\partial a_j^k}.$$

From the definition of  $a_l^{k+1}$

$$a_l^{k+1} = \sum_{j=1}^{r^k} w_{jl}^{k+1} g(a_j^k),$$

where  $g(x)$  is the activation function for the hidden layers,

$$\frac{\partial a_l^{k+1}}{\partial a_j^k} = w_{jl}^{k+1} g'(a_j^k).$$

## Backpropagation -cont'd- :

Plugging this into the latter equation yields a final equation for the error term  $\delta_j^k$  in the hidden layers, called the *backpropagation* formula :

$$\delta_j^k = \sum_{l=1}^{r^{k+1}} \delta_l^{k+1} w_{jl}^{k+1} g'(a_j^k) = g'(a_j^k) \sum_{l=1}^{r^{k+1}} w_{jl}^{k+1} \delta_l^{k+1}.$$

And putting all equations together :

$$\frac{\partial L}{\partial w_{ij}^k} = \delta_j^k o_i^{k-1} = g'(a_j^k) o_i^{k-1} \sum_{l=1}^{r^{k+1}} w_{jl}^{k+1} \delta_l^{k+1}.$$

$\Rightarrow$  gradient values for updating weights at layer  $k$  are computed from the gradient  $\frac{\partial L_d}{\partial a_l^{k+1}}$  used for updating layer  $k+1$

### Computation principle

For each pair  $(x_d^i, y_d^i)$ , compute the output of each neuron by going forward long the network. Then , during a of back propagation of the errors, update all the weights going from the last hidden layer toward the first one.

## BACKPROPAGATION ALGO, Main equations

For the partial derivatives,

$$\frac{\partial L_d}{\partial w_{ij}^k} = \delta_j^k o_i^{k-1}$$

where, for the final layer's error term,

$$\delta_1^m = g'_o(a_1^m) L'(f(x_d), y_d)$$

where, for the hidden layers' error terms,

$$\delta_j^k = g'(a_j^k) \sum_{l=1}^{r^{k+1}} w_{jl}^{k+1} \delta_l^{k+1}.$$

For combining the partial derivatives for each input-output pair,

$$\frac{\partial L(X, \theta)}{\partial w_{ij}^k} = \frac{1}{N} \sum_{d=1}^N \frac{\partial L_d}{\partial w_{ij}^k}.$$

For updating the weights,

$$\Delta w_{ij}^k = -\nu \frac{\partial L(X, \theta)}{\partial w_{ij}^k}.$$

## Backpropagation algorithm summary

- 1) Calculate the **forward phase** for each  $(x_d, y_d)$ ; store the results  $\hat{y}_d$ ,  $a_j^k$ , and  $o_j^k$  for each node  $j$  in layer  $k$  by proceeding from layer 0, to layer  $m$ , the output layer.
- 2) Calculate the **backward phase** for each  $(x_d, y_d)$ ; store  $\frac{\partial L_d}{\partial w_{ij}^k}$  for each weight  $w_{ij}^k$ . Proceed from output layer  $m$ , to layer 1, the input layer.
  - a) Evaluate  $\delta_1^m$
  - b) Backpropagate the error terms for the hidden layers  $\delta_j^k$ , working backwards
  - c) Evaluate the partial derivatives of the individual error  $L_d$  with respect to  $w_{ij}^k$
- 3) Combine the individual gradients for each input-output pair to get the total gradient for the entire set  $X = \{(x_1, y_1), \dots, (x_N, y_N)\}$  (a simple average of the individual gradients).
- 4) Update the weights according to the learning rate  $\alpha$  and total gradient  $\frac{\partial L(X, \theta)}{\partial w_{ij}^k}$

## Remarks

In the classic formulation,  
for hidden nodes ( $g(x) = \sigma(x)$ ) (sigmoidal function)  
and the output activation function is ( $g_o(x) = x$ )

$$g'(x) = \frac{\partial \sigma(x)}{\partial x} = \sigma(x)(1 - \sigma(x)).$$

$$g'_o(x) = \frac{\partial g_o(x)}{\partial x} = \frac{\partial x}{\partial x} = 1.$$

→ No need to remember the activation values  $a_1^m$  and  $a_j^k$  in addition to the output values  $o_1^m$  and  $o_j^k$ , greatly reducing the memory footprint of the algorithm.

**BUT** gradient descent algorithm may be infeasible when the training data size is huge. Thus, a **stochastic** version of the algorithm is often used instead.

## Remarks -cont'd-

- Empirical risk minimization for multilayer perceptron is an **ill-posed** and **ill-conditioned NON CONVEX** problem.
- Gradient values in **the first hidden layers** often takes either **too large** (explosion) or too low values (vanishing gradient, leading to slow down learning convergence).
- Weights initialization values, learning rate, choice of  $g()$ ,  $m$ ,  $r_m$  do all influence the result! first results date back to Hinton, 2006
- To **avoid saturation effects** of node outputs (either to 0 or to 1),  $l_2$  regularization of  $L(f(x), y)$  may be applied on  $\theta$ .



## Stochastic Gradient method

At each iteration, rather than computing

$$\nabla_{\theta} L(\mathbf{X}) = \nabla_{\theta} \left( \sum_{d=1}^N L(x_d) \right) = \sum_{d=1}^N \nabla_{\theta} L(x_d)$$

stochastic gradient descent randomly samples  $d$  at uniform and computes  $\nabla_{\theta} L(x_d)$  instead :

SGD uses  $\nabla L(x_d)$  as an unbiased estimator of  $\nabla L(\mathbf{X})$  ;


$$\mathbb{E} [\nabla_{\theta} L(x_d)] = \mathbb{E} \left[ \frac{1}{k} \sum_{i=1}^k \nabla L(x_k) \right] = \nabla L(\mathbf{X}).$$

In a generalized case, at each iteration a mini-batch  $\mathcal{B}$  that consists of indices for training data instances may be sampled at uniform with replacement.

$$\nabla L_{\mathcal{B}}(\mathbf{X}) = \frac{1}{|\mathcal{B}|} \sum_{d \in \mathcal{B}} \nabla L(x_d)$$

update  $\theta$  as

$$\theta := \theta - \eta \nabla L_{\mathcal{B}}(\mathbf{X})$$



- ▶ The per-iteration computational cost is  $\mathcal{O}(|\mathcal{B}|)$ . Thus, when the mini-batch size is small, the computational cost at each iteration is light.
- ▶ If the training data set has many redundant data instances, stochastic gradients may be so close to the true gradient  $\nabla_{\theta} L(\mathbf{X})$  that a small number of iterations will find useful solutions to the optimization problem.
- ▶ Stochastic gradient descent can be considered as offering a regularization effect especially when the mini-batch size is small due to the randomness and noise in the mini-batch sampling.
- ▶ Certain hardware processes mini-batches of specific sizes more efficiently.

[See MLPC\\_sonar\\_example.ipynb](#)

## Backprop in Practice

Y LeCun

- Use ReLU non-linearities
- Use cross-entropy loss for classification
- Use Stochastic Gradient Descent on minibatches
- Shuffle the training samples (← very important)
- Normalize the input variables (zero mean, unit variance)
- Schedule to decrease the learning rate
- Use a bit of L1 or L2 regularization on the weights (or a combination)
  - ▶ But it's best to turn it on after a couple of epochs
- Use "dropout" for regularization
- Lots more in [LeCun et al. "Efficient Backprop" 1998]
- Lots, lots more in "Neural Networks, Tricks of the Trade" (2012 edition) edited by G. Montavon, G. B. Orr, and K-R Müller (Springer)
- More recent: Deep Learning (MIT Press book in preparation)


Save

## AutoEncoder

Images for this section were adapted from <https://www.jeremyjordan.me/autoencoders/>

- Autoencoders are **Unsupervised** Neural Networks, designed for **Representation learning** and/or dimension reduction.
- main idea : impose a **bottleneck** in the network to **compressed** knowledge representation of the input.

*rk* : This assumes that the data are structured (input features are correlated) as for e.g. iid data, such compression will be very difficult if not impossible without losing much information.




## AutoEncoder principle

⇒ Formulate the problem as a **supervised learning** problem whose output is  $\{x^i\}$

⇒ The empirical risk to minimize is thus  $L(f(x), x)$  : the bottleneck plays a key role (otherwise the network simply passes the values to the output).

⇒ if linear activation function were used, that would perform PCA like dimension reduction

Linear vs nonlinear dimensionality reduction



The AutoEncoder must be :


- ▶ sensitive enough to the inputs, to built an accurate reconstruction
- ▶ insensitive enough to the inputs to avoid overfitting

This requires to regularize the loss function of the form

$$L(f(x), x) + \textit{regularization}$$

## UnderComplete AE

Limit of the flow of information going through the NN by limiting the nb of nodes in the hidden layers :



No explicit regularization term is required here.

If **too many nodes**, i.e. high capacity -in the sense of Vapnik-, the AE may be capable of learning a way to **simply memorize the data**. The primary aim to discover latent variable cannot be attained !

## Sparse AE

**Idea** is to keep the number of nodes in hidden layers quite large, but regularize the loss function by penalizing activations within a layer ( $\neq$  weights regularisation)  $\Rightarrow$  only a small nb of neurons are activated.

$l_1$  regularization

For layer  $k$  :

$$(f(x), x) + \lambda \sum_{j=1}^{r_k} |o_j^k|$$

**Remind** that in the notation of the previous section, activation of a neuron was noted  $o_j^k = g(a_j^k)$ . Activations are in  $[0, 1]$  or  $[-1, 1]$ , depending on the choice of  $g()$ . Below, we assume  $o_j^k \in [0, 1]$ .

## Sparse AE - cont'd-

### Kullback-Leibler regularization


Let  $\hat{\rho}_j^k \stackrel{def}{=} \frac{1}{n} \sum_{i=1}^n o_j^k(x)$  be the average activation of neuron  $i$  in layer  $k$ , estimated over a collection of  $n$  samples  $\{x^i, i = 1 \dots n\}$ .

$$L(f(x), x) + \lambda \sum_{k=1}^m \sum_{j=1}^{r_k} \text{KL}(\mathcal{B}_\rho || \mathcal{B}_{\hat{\rho}_j^k})$$

where  $\mathcal{B}_\rho$  is the Bernoulli process of parameter  $\rho$ , thus

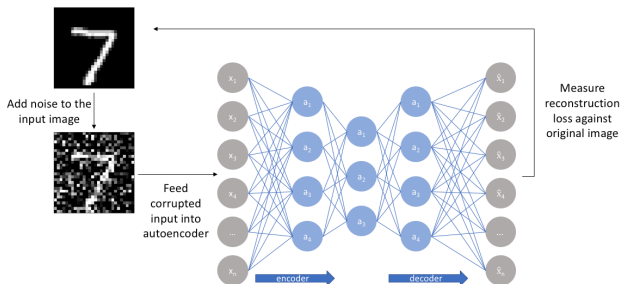
$$\text{KL}(\mathcal{B}_\rho || \mathcal{B}_{\hat{\rho}_j^k}) = \rho \log \frac{\rho}{\hat{\rho}_j^k} + (1 - \rho) \log \frac{1 - \rho}{1 - \hat{\rho}_j^k}$$

$\Rightarrow \rho$  acts as a "sparsity parameter";  
small values of  $\rho$  correspond to low  
probability for the neuron to fire.





## AE for denoising : principle



### Alternate approach for denoising

Force the activation of the hidden layer to be weakly sensitive to small deviations of the inputs

$$L(f(x), x) + \lambda \sum_{k=1}^m \sum_{j=1}^{r_k} \|\nabla_x o_j^k(x)\|^2$$